

LangGraph (Persistence) — Full English Note Summary

1. Introduction — Why “Persistence” Matters

- **Topic:** *Persistence in LangGraph*
 - **Core idea:** It's the ability to save and restore the state of a workflow over time.
 - **Why important:**
 - It's a **foundational concept**; many advanced features (like memory, fault-tolerance, human-in-loop) are built on it.
 - Without persistence, once a workflow finishes, all its data in memory is lost.
-

2. Revisiting Key LangGraph Concepts

Before understanding persistence, recall:

1. **Graph concept** — You decompose a goal into multiple tasks (nodes) connected by edges defining execution order.
2. **State concept** —
 - State = a Python **dictionary** holding workflow data (`messages` , `results` , etc.).
 - Every node can **read & write** state values.
 - Example: Chatbot state stores exchanged messages.

➡ Normally, after a workflow ends, all these state values vanish from RAM.

3. What is Persistence?

Definition:

“Persistence in LangGraph refers to the ability to save and restore the state of a workflow over time.”

- By default, state is temporary — erased after execution.
 - **With persistence:** You can **store** the state (final + intermediate values) in a **database** (or in-memory store).
 - Later, you can **retrieve** or **resume** a workflow from where it stopped.
- ✅ So persistence makes workflows *fault-tolerant, resumable, and inspectable*.
-

4. Key Specialty — Stores *All* States, Not Just Final

- Persistence doesn't only store the **final** state; it also stores **intermediate states** after every node or step.
 - This means:
 - You can see the state's evolution ($A \rightarrow B \rightarrow C$).
 - If a workflow crashes midway, you can **resume from the last saved point**, not from the beginning.
 - This enables **fault tolerance** and **progress tracking**.
-

5. How Persistence Works Internally

5.1 The "Checkpoint" System

- LangGraph implements persistence through **Checkpointer**s.
- A *Checkpointer* breaks the workflow execution into **supersteps**, placing a **checkpoint** after each.
- At every checkpoint, **current state values** (intermediate + final) are saved.

Superstep Example:

`Start` → Node1 → {Node2, Node3 in parallel} → `End`

→ Has 3 supersteps → 3 checkpoints created

→ State saved at each checkpoint.

5.2 Types of Checkpointers

- **InMemorySaver** — stores states in RAM (demo use).
 - **PostgresSaver**, **RedisSaver** — production checkpointers that persist states in databases.
-

6. Thread Concept (Multiple Executions)

- Each workflow execution gets a unique **thread_id**.
- Every state (and checkpoint) is stored **against its thread_id**.
- Benefits:
 - You can resume or fetch states for a **specific run**.
 - Enables multiple concurrent or past runs of the same workflow.

Example:

`Thread 1` → `topic: "Pizza"` → `Joke workflow`

`Thread 2` → `topic: "Pasta"` → `Joke workflow`

→ Both persisted separately; retrievable anytime.



7. Code Example (Sequential Workflow with Persistence)

Task: Generate a joke and explanation using an LLM.

Workflow:

START



generate_joke



generate_explanation



END

State

topic: `str` joke: `str` explanation: `str`

Steps

1. Import libraries

```
from langgraph.checkpoint.memory import InMemorySaver
```

2. **Build nodes** (generate_joke , generate_explanation) — each writes partial state.

3. **Add edges:** start → joke → explanation → end.

4. Create checkpointer & compile

```
checker = InMemorySaver() graph = workflow.compile(checkpointer=checker)
```

5. Run with thread_id

```
config = {"configurable": {"thread_id": "1"}} graph.invoke({"topic": "pizza"},  
config=config)
```

6. Retrieve states

```
graph.get_state(config=config) # final state graph.get_state_history(config=config) #  
all checkpoints
```

✅ Each checkpoint stores a snapshot of state values (empty → after topic → after joke → after explanation).

8. Advanced Scenarios Enabled by Persistence

8.1 Fault Tolerance

- If the workflow crashes midway (server down, API error, etc.), it can **resume from the last saved checkpoint** instead of restarting.

8.2 Short-Term Memory (Chatbots)

- Essential for **resume chat** functionality (like ChatGPT).
- Old messages are stored via persistence; user can continue an old chat later.





8.3 Human-in-the-Loop (HITL)

- During execution, LangGraph can **pause** at a node waiting for user input.
- Persistence ensures the system can **resume later** from that pause point.

8.4 Time Travel (Debugging / Replay)

- You can **replay** or **branch** a workflow from any past checkpoint.
- Helps in debugging or re-running sections with different inputs (e.g., change "Pizza" → "Samosa" topic).

9. Practical Benefits of Persistence

Feature	Description	Example
 Short-Term Memory	Save & recall past conversation context	Chatbot resume
 Fault Tolerance	Resume after crash from last saved state	Long workflows
 Human-in-the-Loop	Pause for manual input & resume later	Approval gate before posting
 Time Travel	Replay / branch from older checkpoints	Debug or regenerate results

10. Summary — Key Implementation Points

1. Always use **checkpointer** in `workflow.compile()` .
2. Provide a **thread_id** when invoking (`configurable={"thread_id": id}`).
3. Use `get_state()` for final results, `get_state_history()` for checkpoint timeline.

4. For production, replace `InMemorySaver` with database-backed checkpointer (Postgres, Redis).
 5. You can **resume**, **replay**, or **update state** at any checkpoint.
-

💡 11. Conceptual Map (Cheat-Sheet)

